

gowkhtmltopdf - Architecture

A work-alike for *wkhtmltopdf* built for controlled reports: invoices, statements, tables, multi-page documents with headers/footers, TOCs and PDF outlines. Everything from **load** to **write** runs inside the Go binary.

DOCS

11 deep-dives in documentation/architecture/

DOMAINS

10 sections on this page

SCALE

~**71 600** lines of Go

COMMIT


ef526f9 2026-08-11

INPUTS

SHARED ENGINE

OUTPUTS

Conversion pipeline one document, seven stages



Reading order jump to any domain

1. [01 · Entrypoints & CLI](#)

2. [02 · Public library API](#)

3. [03 · Settings system](#)

4. [04 · Load layer](#)

5. [05 · HTML parser](#)

6. [06 · CSS subsystem](#)

7. [07 · Layout engine](#)
8. [08 · Convert pipeline](#)

9. [09 · PDF writer](#)

10. [10 · Image output & SVG](#)

11. [Dependency DAG](#)

12. [PDF vs image](#)

13. [Security](#)

01

Entrypoints & CLI

Three entry points: two CLIs and the library; every flag funnels through dotted settings into the `convert.Request` seam.

The outermost shell turns raw process arguments into a settings-shaped `cli.Command` and hands it to the application adapters that build engine requests. One parser serves both binaries through a `Mode` bitmask, and the parser is the only place that understands the *wkhtmltopdf*-compatible multi-object grammar: `page`, `cover` and `toc` objects plus free positionals.

ONE PARSER, TWO BINARIES - ARGV IN, REQUEST OUT

```
$ gowkhtmltopdf --page-size A4 cover cover.html toc page ch1.html out.pdf
write PDF: 3 pages · outline depth 2 · exit 0
```

CODE	MEANING	SOURCE
0	ok	success

CODE	MEANING	SOURCE
1	error	parse or engine failure
2	HTTP 404	settings.HttpStatusError
3	HTTP 401	settings.HttpStatusError

Flags land in settings through dotted names: `--page-size` writes `size.pagesize`; free positionals become `Command.Objects` and the last one becomes `Output`. Errors surface as `gowkhtmltopdf: unknown option --dpi` on stderr with exit code 1, while HTTP failures report `2` or `3`. Long option names may split at word boundaries when printed.



`cmd/gowkhtmltopdf` PDF binary main - parse (ModePDF), doc-flag dispatch, signals, exit code; 83 thin lines
`cmd/gowkhtmltoimage` Image binary main - parse (ModelImage), no TOC-XSL dump, exit code; 53 thin lines
`internal/cli` Parser core: Command, Parse, Mode bitmask, ~71 registered flags, ExitCode mapping
`internal/app` Adapters BuildPDFRequest / RunImage turn Command into engine requests

- `cli.Command` - parse result: Global/Image settings, Objects, Output / OutputWriter
- `cli.Parse(argv, Mode)` - token loop, object grammar, pending remapping
- `cli.ExitCode(err)` - 0/1/2/3 via `HttpErrorCode`
- `settings.Set("size.pagesize", ...)` - the dotted settings home appliers write
- `newFreshObject()` - TOC objects never consume pending - no ghost page

Security: local-file ACL is opt-in (`--enable-local-file-access` dual-writes `load.blocklocalfileaccess`); Policy A rejects inert wkhtmltopdf stubs (`--dpi`, `--javascript-delay`) as unknown rather than no-ops; validation precedes output truncation; the CLI spawns no subprocess and cancels cleanly on SIGINT/SIGTERM via a context.

Position: `argv` → `settings` → `Request` (left edge of the pipeline) · see `documentation/architecture/01-entrpoints-cli.md`

02 Public library API

Converter / ImageConverter public API: typed and dotted dialects, deep-copy ownership, in-memory output.

The root package `gowkhtmltopdf` is the idiomatic, stdlib-only Go surface for the engine. It configures a conversion with either wkhtmltopdf-compatible dotted strings or compile-time-discoverable typed builders, and drives the pipeline: `Converter.Convert(ctx)` reduces to an internal `convert.Request` handed to `convert.Run` (PDF) or `imageout.RunRequest` (image), with all output captured in memory. This package also the module boundary that enforces the product rule - `pure Go, no cgo, stdlib + internal only` - so any new external dependency would have to cross its import list first.

`GlobalSettings.Set("size.pagesize", "A4")`
dotted dialect

`PdfGlobalOptions / PDFRequest`
typed dialect

`ImageConverter / RunImage`
image mode

R
O
O
T
P
A
C

convert.Request

one neutral seam - the engine never parses argv

render.Run →

PDF

imageout.RunRequest → PNG

/ JPEG

fig. 2 - two dialects, one seam. `api.go` (960 lines) never imports `internal/cli`; both frontends meet at the shared request union (2026-08-11, commit ef526f9).

`api.go`

`api_test.go`

`doc.go`

`internal/convert,`
`settings, imageout,`
`line, errs`

all exported types: settings wrappers, converters, typed requests, hooks, sentinels

round trips, deep-copy mutation isolation, cancellation, callbacks, PNG/JPEG decode

quick start, dotted settings, ACL pair, in-memory source kind, thread-safety contract

the five internal dependencies the root package adapts

- `Converter` + `ImageConverter` - lifecycle: `New`, `AddObject`, `Convert(ctx)`, `Output()`
- `GlobalSettings` / `ObjectSettings` - dotted `Set` / `Get` via the reflect key tables
- `PDFRequest` / `ImageRequest` + `RunPDF` / `RunImage` - typed one-shot; `Now func() time.Time` gives byte-stable output
- `ConvertHTML` - one-shot helper: clone global, `SetBody(html, base)`, convert, output
- `Version()` - banner `0.12.7-dev (gowkhtmltopdf pure-go)`, tracking `LibraryVersion 0.12.7-dev`
- Sentinels - `ErrNoPageObjectsAdded`, `ErrEmptyHTML`, `ErrNilContext`, matchable with `errors.Is`
- `convertHooks` + `lineLog` - engine log lines classified by `line.SeverityOf` into `OnInfo` / `OnWarn` / `OnError`

Surface	Contract
<code>Converter</code>	PDF, multi-object, <code>OnPhase</code> / <code>OnProgress</code> (final 100); no <code>Now</code> hook
<code>ImageConverter</code>	PNG/JPEG, single most-recent page, lazy zero-value init, no phase callbacks
<code>RunPDF</code> / <code>RunImage</code>	typed path - <code>Now</code> injectable for reproducible metadata

Security posture: local-file ACL is default-deny and needs the same two-step opt-in as the CLI (`enablelocalfileaccess=true + load.blocklocalfileaccess=false`); `SetBody` / `ConvertHTML` bypass URL guessing so in-memory HTML cannot be coerced into SSRF; context cancellation threads through every load. Limits: no `HttpErrorCode()` on `Converter` (CLI-only today), no TOC as a first-class object, one image page only. Every boundary copies: `SetBody`, `AddObject`, `Output`.

Same seam as the CLIs: `api.go` → `convertHooks` → `Request` → `render.Run` - deep-dive:

`documentation/architecture/02-library-api.md`

03 Settings system & errors

wkhtmltopdf-style settings model with reflection key tables; settings and errs are leaf packages.

`internal/settings` is the single source of truth for every user-facing knob - page geometry, margins, headers/footers, TOC behaviour, load policy, image output - and the only package that speaks the `wkhtmltopdf`

dotted-name vocabulary (`margin.top`, `load.jsdelay`, `web.background`). It sits upstream of everything: argv and library calls funnel through `Set("margin.top", "25mm")`, and every engine package reads the typed structs it default-initialises. Its companion leaf `internal/errs` owns the shared sentinel errors across the library boundary.

DOTTED-KEY PATH - ONE KNOB, TYPED ALL THE WAY DOWN

`--margin-top 25mm` → `Set("margin.top", "25mm")` → `Margin.Top = 25 mm`

`10 mm` → `28.346 pt`

Booleans coerce liberally (`"` / `true` / `1` / `yes` / `on`); unknown keys fail loudly - a typo like `margin.topx` errors instead of degrading.

POLICY A SETTINGS HONESTY

Only options with an engine consumer get typed fields. Inert `wkhtml` keys (`dpi`, `javascript`, `js-delay`, `produce-forms`, `default-encoding`, `.`) sink into an `Ignored` map: *accepted* so existing invocations do not hard-fail, *inert* so they are never silently promoted to typed stubs. Origin: the 2026-08-06 ponytail review, phase 1.


accepted-inert `wkhtml`
keys, round-trippable via
`Get`


ISO/ANSI page sizes -
A4 = 595.28 × 841.89 pt

`internal/settings/settings.go`

Typed model: `PdfGlobal`, `PdfObject`, `ImageGlobal` plus `Web`, `LoadGlobal`, `LoadPage`, `HeaderFooter`, `TableOfContent` - `wkhtmltopdf`-compatible defaults (A4 portrait, 10 mm margins, depth-4 outline).

`internal/settings/reflect.go`

Descriptor engine: dotted `Set` / `Get`, type-coercing setters, ignored-key tables - built once at init, zero per-call reflection cost.

`internal/settings/unitreal.go`

`UnitReal` (`10mm`, `1.5in`, `12pt`, `96px`, `100%`), the 23-entry page-size table, and

`· pagesize.go · httperror.go`

`HttpStatusError` (`404` → `2`, `401` → `3`, else `1`).

`internal/errs/errs.go`

Shared sentinels: `ErrNilContext`, `ErrNilLoader`, `ErrNilCommand`, `ErrNilRequest` - re-exported by `api.go` and `internal/app`.

- `PdfGlobal` / `PdfObject` / `ImageGlobal` - the three settings payloads carried by `cli.Command` and `convert.Request`.
- `Set` / `Get` - dotted-key surface; accepted inert keys round-trip, unknown keys error.
- `UnitReal` - unit-suffix scalar; `Points()` / `Mm()` conversion.
- `ParsePageSize` — case-insensitive ISO/ANSI names to points; empty string → A4.
- `HttpErrorCode` — `404` → `2`, `401` → `3`, else `1`; consumed by `cli.ExitCode` via an interface check.

SECURITY & LIMITS

Local file access is denied by default (`BlockLocalFileAccess: true`); `--allow` prefixes are the only ACL widening;

proxies must be absolute `http(s)` URLs. Inert-key acceptance is a compat surface, not an attack surface - the engine never executes scripts or creates forms.


Known gaps: `default-encoding` is accepted-ignored; `size.pagesize` dual-write is transitional; object-level `web.background` is inert; the typed builder covers global PDF options only.

Position: upstream of every entry point — dotted Set/Get → UnitReal → page sizes → convert pageGeometry. Source: [documentation/architecture/03-settings.md](#) (commit ef526f9, 2026-08-11).

04 Load layer - URLs, I/O & ACL

Resource-fetching seam and trust boundary: URL guessing, local-file ACL, HTTP hardening, size caps.

`internal/load` owns the first pipeline stage and the two resource seams the engine depends on: the **primary load** (one `Loader.Load` per body object, cover, TOC, or header/footer) and the **subresource seam** (`ResourceContext.Fetch` for linked CSS, images and `@font-face` URLs). It is the lowest internal package - only `internal/settings` and the Go standard library sit beneath it - and it carries the product security posture: default-deny local-file ACL, connect/response timeouts, a 10-redirect cap, a 100 MiB body cap, and a strict UTF-8/ASCII charset gate.



Primary load

`Loader.Load(ctx, input, pageLoad)` resolves pages, covers, TOCs and header/footer HTML like top-level documents. Budgets: 30 s connect, 60 s response default, 100 MiB cap, max 10 redirects, per-page `--timeout` override.

Subresource seam

`ResourceContext.Fetch(ref)` binds a document base plus the per-page policy, so CSS, images and fonts inherit the same ACL and caps - e.g. `internal/convert/prepare/styles.go` collects linked stylesheets through this one narrow seam.

[internal/load/load.go](#) URL guessing, HTTP/file/data: fetching, cookie jar, proxy, auth, POST, ACL, caps, charset gate (1,322 lines)

[internal/load/load_test.go](#) 30 behavioural tests: ACL matrix, body caps, timeouts, redirects, subresources, concurrency

[internal/load/doc.go](#) Package doc: "reimplements the MultiPageLoader orchestration layer"

- `Kind` - resolved input class: HTTP | File | Inline
- `Resource` - fetched document: URL, Base, Body, Skip
- `ResourceContext` - the one sanctioned subresource seam
- `AccessController` - default-deny ACL, symlink-resolved prefixes
- `NewLoaderWithError` - fail-fast constructor used at the request boundary

Global enable	Object block	Read allowed
off	on (default)	no

Global enable	Object block	Read allowed
off	off	no
on	on	no
on	off	yes

URL shapes handled: `http://host/path`, `file://localhost/tmp/report.html` (remote file hosts refused), `data:` URIs, and bare input defaulting to `http://<input>` - a direct port of `wkhtmltopdf`'s `guessUrlFromString`.

Trust boundary owner: every input string crosses here into raw bytes; nothing below `internal/load` touches the network or filesystem.

Security: default-deny ACL with `--allow` prefix expansion; no JavaScript ever; oversized bodies are rejected, never truncated; the TOCTOU window between ACL check and read is accepted residual risk; HTTP egress to `localhost/RFC1918` is intentionally unrestricted (input-trust model). Full model: `documentation/THREAT-MODEL.md`.

First pipeline stage feeding `internal/html` + the `prepare` subresource seam · deep-dive: `documentation/architecture/04-load.md`

05

HTML parser: tokenizer & tree

Generic tokenizer and tree builder: entities, tolerant parsing, script content stripped at layout.

`internal/html` is the parse stage of the pipeline: a hand-rolled, *stdlib-only* leaf that owns the **only DOM in the repo** - every downstream stage walks its `Node` tree. Scanning is streaming: `scanTokens` emits tokens through a sink callback while the tree is built incrementally in the same pass, so a whole-token slice is never materialized.

Token constructs the scanner recognizes

text run

<!-- comment -->

<!doctype>

<?PI?>
skipped

</end>

<start>

<raw>
verbatim

Raw-text elements - `script`, `style`, `textarea`, `title` - are captured verbatim up to their real closing tag, then hidden by the UA sheet (`display:none`, `style_values.go:986`): script content is stripped at *layout*, not at parse. Entities decode once, at text-append and attribute application, via `UnescapeEntities` over `stdlib.html.UnescapeString`.

A ruby annotation parses as one inline run - base, parens and `reading` are ordinary inline nodes (CJK ruby needs `--font-path` for glyphs):

XML([Extensible Markup Language](#)) 2026-08-11 · `html.go` ≈ 832 lines · `allowlistedbyconstruction`

Tolerance is a feature: browser-grade-recovery is deliberately out of scope, while malformed input degrades to a usable tree- six sentinel errors, never a panic.

[internal/html/html.go](#)

Tokenizer, tree builder, Node model, walkers, element vocabularies (~832 lines)

[internal/html/entities.go](#)

UnescapeEntities - &-gated stdlib entity decoding

[internal/html/html_test.go](#)

Same-package tests: tokenizer, tree, malformed-input table (~901 lines)

Consumers

css.Match, layout style resolution, outline heading collection, convert HF/TOC/islands

- `Parse` / `ParseDocument` - string vs bytes (+BOM) entry points
- `Node.Walk` - pre-order traversal used by every consumer
- `scanTokens` - streaming scanner with sink callback
- `autoClose` - implicit-close vocabulary (li, p, tr, td, th.)
- `mergeRootElement` - html/head/body merge instead of nesting
- `uaDec1s` - UA sheet hiding script/style/title/head/meta/link

Security: no script execution by construction - event-handler attributes are inert data; charset is enforced (UTF-8/ASCII only) at the load seam before parsing; hostile input cannot panic.

Position: load - [html](#) - (css matching, layout) - deep-dive: [documentation/architecture/05-html-parser.md](#)

06

CSS subsystem: parse, selectors & cascade

CSS subset: selectors, cascade, media and container queries; unknown rules degrade; `:has` and `:target` supported.

`internal/css` is a pure parsing, matching and value library that sits between the HTML tree and the layout engine. It owns the *CSS subset* the converter accepts and reports matches and specificity to `internal/layout`, which performs the cascade. It never fetches resources, never resolves page size, and never decides layout - the contract is quoted in its package doc (*css.go:1-18*).

parse-time

match-time

value helpers

never panics

html

css

layout

pdf / image

Supported surface: `type`, `.class`, `#id`, attribute selectors, `:first-child` / `:last-child` / `:nth-child`, `:has()`, `:not()`, descendant/child/sibling combinators, `@media`, `@container`, `!important`, and inline `style=""`. Everything else degrades without panicking- and unsupported selectors never silently degrade to the host element (`writePseudoLiteral` keeps unknown pseudos on the compound, so `li:target` can never paint every list item).

Degrade rules - how the subset stays honest

- `garbage` is dropped: only unbalanced braces error, and `ParseNeverPanics` pins the no-panic property on adversarial input.
- `never-match` pseudos (`hover`, `active`, `focus`, `target`) stay on the compound and match `false` - a static PDF has no pointer state.
- `skip` at-rules: `@import`, `@supports`, `@layer`, `@keyframes`, nesting are silently skipped.
- `parse-early`: media preludes stay raw strings (the viewport is only known at conversion); `:nth-child` arguments compile to integer form so matching is pure arithmetic.

Known gaps - what degrades today

WIRE FORMAT - ONE DECLARATION BLOCK

```
--accent : #b4532a ; custom property
color : var(--accent, #b4532a) ; var() with fallback
font-size : 1.1rem ; !important single global tier
```

Unit math is tuned for IEEE float cancellation: 25.4mm is computed as $\text{val} \times 72 \div 25.4$ so it cancels cleanly to 72pt instead of 71.999 ; px maps at 0.75 ($96\text{css px}/\text{in} \rightarrow 72\text{pt}/\text{in}$); the media em

base is 12pt . Specificity is Selectors-4: `:has()` / `:not()` contribute the max of their arguments - a triple (a, b, c) ¹ with a = ids, b = classes/attrs/pseudos, c = types.

¹ Functional pseudos contribute the specificity of their most specific argument, not a flat class-level count.

<code>internal/css/css.go</code>	Sheet/rule/selector model, parser, at-rules, right-to-left matching, specificity (≈ 1784 lines)
<code>internal/css/values.go</code>	Inline style, <code>!important</code> , length/color/font parsing, custom-property resolution
<code>internal/css/container.go</code>	<code>@container</code> prelude, boolean condition tree, <code>LengthToPt</code>
<code>internal/css/has.go</code>	<code>:has()</code> / <code>:not()</code> combinator walk; <code>@media</code> evaluation
<code>media.go</code>	

- `css.Parse` - whole-sheet parse; only unbalanced braces error
- `css.Match` / `MatchPseudo` - selector vs node, right-to-left; `::before/::after` shapes
- `css.Specificity` - (a, b, c) triple; `:has()/:not()` max-of-arguments
- `css.MediaMatches` - raw `@media` prelude vs conversion viewport
- `css.ResolveCustomProps` - the single place custom-property policy lives
- `css.HasContainerRules` - fast gate; layout skips its second style pass when false
- `css.ParseSelectors` - strict list parse for `--exclude-from-outline`

Import rule: `internal/css` → `internal/html` → `stdlib` only; `css` never imports `layout`, `convert`, `pdf`, `load`, or `settings` - `viewport/media/container` context is always passed as arguments. **Security:** `css` executes no code and fetches nothing; `url()` is honored only in `@font-face src` via the loader's ACL, and stylesheet rule budgets are capped in preparation (soft warn 25 000, hard cap 1 000 000).

Position: `html` - `css` - layout style resolution · measured 2026-08 · deep-dive: `documentation/architecture/06-css.md`

07

Layout engine & line breaking

Style cascade plus all formatting contexts, line breaking, pagination, emitting the display list.

`internal/layout` is the largest domain in the repo: it runs the cascade and builds an absolute-positioned *display list* of drawing ops on a continuous y-down canvas. Layout also owns pagination - page-break policies, thread repetition, sticky clamps, and the op-split machinery. One engine feeds three painters: PDF, header/footer bands, and the PNG/JPEG rasterizer.



layout and paginate both live in internal/layout; the same ops feed pdf.Paint, PaintBand (HTML header/footer) and the imageout rasterizer.

One display-list Op OP KINDS

Every op carries a stable ID. Kinds: OpFillRect, OpStrokeRect, OpLine, OpText, OpImage, OpLinkURI, OpBullet. Rects that split across a page boundary keep the source ID, so element-to-op ownership survives fragmentation.

Final z-order is resolved once by PaintOrder and shared by all three backends.

Formatting contexts (report-lite subset)

- block & inline flow
- flex lite (row / column)
- multicol lite
- relative / absolute / fixed lite
- tables: colspan, rowspan, thead repeat
- grid lite + subgrid-inherit
- float lite + clear
- print sticky, 2D transforms

Key entry points (layout)

- LayoutContext / WithWorkspace - the layout passes
- PaintContext - paginate and paint into a pdf.Document
- PaintBandContext - single clipped band for HTML header/footer
- CloneResult - deep copy for the TOC page-count fixpoint

Scale: ~33.6k lines incl. tests, 27 production files, 62 test files, 244 test functions.

internal/layout

cascade, boxes, ops, pagination, paint; entry points LayoutContext / PaintContext

internal/line

naming note: log-line severity protocol (line.Emit), NOT text breaking - line breaking lives in layout inline.go / inline_collect.go

internal/svg

Rasterize is used for inline SVG images; layout stays one layer below convert and never imports convert/load/settings


- Result - Ops, Pages []int, Locations []ElementLocation
- ResolvedStyle - ~1.3 KB per style, interned per run via styleStore
- ElementLocation - feeds outline.Lookup and #frag link assembly
- Options - Images func(src) callback keeps the ACL-gated loader injected
- Workspace.Release - display-list storage reuse for the page-islands path

Images arrive only through the injected Images callback (ACL-gated); layout never opens URLs and never runs JS. Work is bounded: the pagination fixpoint caps at 10 iterations, and SVG rasterizes at a 1024 px cap. Lite limits: positioning, sticky, float, flex/grid/multicol are report subsets; overflow is not general clipping; no vertical writing modes; HTML header/footer bands clip rather than paginate.

style - boxes - paginate - display list (consumed by pdf OR imageout) · deep-dive: [documentation/architecture/07-layout.md](#)

Orchestration hub: the Request union, 3-stage render.Pipeline, headers/footers, TOC fixpoint, outline, page islands.

internal/convert is the only package that coordinates a complete conversion. It owns the neutral `Request` union (convert.go:51), runs the per-object render loop, and drives every document-wide pass - TOC, outline, links, headers/footers, copies - before writing a single PDF. The lifecycle is delegated to a three-stage `render.Pipeline`, keeping PDF specifics out of the stage driver.



Lifecycle adapter: `render.Pipeline` (render/pipeline.go) owns stage order and cancellation; `pdfPipeline` (pdf_pipeline.go) owns the PDF detail. Cancellation is checked at every object boundary and between stages.

Assemble, in order

1. `assembleTOC` - two-pass fixed point (`renderTOCObjects`, toc.go:221): `iteration 1` measures with `tocGuess 0`, `iteration 2` rennumbers with the measured total and keeps the final layouts; pages reorder so the TOC comes first.
2. `assembleOutline` - pure heading tree (`internal/outline`, no PDF types) → `emitOutline` → `pdf.Outline` with final page refs.
3. `assembleLinks` - TOC forward/back links plus `#id GoTo` annotations (`applyInternalLinks`).
4. `assembleDocument` - page plan, title, compression, grayscale, creation time.
5. `assembleCopies` - collate / non-collate materialization via `render.Plan`.
6. `assembleHeadersFooters` - final pass; needs [topage] and section/subsection, which only exist once the whole document is final.

Fact: headers/footers are painted once at the very end - after copies - because `[topage]` and section/subsection are only known when the whole document exists. Auto margins reserve the bands at layout time (`effectiveMargins`, hf.go:596); the draw pass clips ink into those same bands.

<code>internal/convert</code>	job model, per-object loop, document-wide passes
<code>internal/convert/prepare</code>	shared load → parse → sheets → fonts phase (same seam used by image mode)
<code>internal/convert/render</code>	Pipeline lifecycle driver + Plan page-index model
<code>internal/outline</code>	pure heading tree, dump XML, section lookup (no layout/PDF result types)

- `Request` - neutral pipeline input: Global, Image, Objects, Now, Output, OutlineOutput (convert.go:51)
- `render.Pipeline` - RenderObjects / Assemble / Finalize (render/pipeline.go:12)
- `runContext` - one conversion lifecycle: loader, fonts, single pdf.Document
- `objectState` - per-object geometry, HF, headings, navigation projection
- `pagePlan` - owner map + Remap after TOC reorder / copies (page_plan.go)
- `drawHeadersFootersResult` - final HF pass with real page numbers (hf.go:637)

Security: the loader (proxy/ACL boundary) is built at the top of `Run` before fonts or layout state; every subresource fetch inherits its ACL, timeouts and body caps. Header/footer values that look like HTML are treated as raw markup, never as URLs - blocking a local-path injection vector. Output sinks are caller-supplied writers; convert never opens paths.

Budgets: 10 000 objects, 1 000 copies, 100 000 pages.

Limits: `--xsl-style-sheet` is not implemented (built-in TOC template, warning + fallback); the TOC fixed point is bounded to two iterations; HTML headers/footers are single-band clamps.

Progress is a (phase, percent) callback:  100 = progressComplete at Finalize; the CLI prints phase lines unless `--quiet`.

Position: Request → RenderObjects → Assemble → Finalize (drives every stage) · deep-dive: 08-convert-pipeline.md

09

PDF 1.4 writer


PDF 1.4 sink: Flate streams, font subsetting, Type0/CID, xref, outlines, byte-stable output.

`internal/pdf` is the lowermost writer layer: it turns the layout display list into a viewer-openable PDF 1.4 file

using only the Go standard library plus one allowlisted shaping module (`go-text/typesetting`). It owns the document object model and serialization, the content-stream operator language, the font pipeline (parse → subset → simple or Type0/CID embed), image embedding under hard caps, and outline/annotation wiring. Image mode reuses its fonts, shaping and glyph outlines but rasterizes instead of writing PDF bytes.

display list → paint → content streams → xref

Font pipeline: untrusted TTF / WOFF1 in, one subset per font per document out



Content-stream language

Every operator is emitted by the `Content` builder: graphics state `q / Q`, paths `m / l / re / f / s`, transforms `cm`, and text `BT / ET / Tf / Tj`, with fonts and images registered into per-page `/Resources` at finalize. Text emits through `TextShow`: pure ASCII fast path, otherwise shape then split mixed Latin / CJK runs.

Serialization

`writeTo` streams objects through a `countingWriter` so every xref offset is exact; all `/FlateDecode` streams use RFC 1950 zlib (raw DEFLATE made pages render empty). A typical tail:

```
0000000000 65535 f
0000000015 00000 n
0000000091 00000 n
trailer << /Size 3 /Root 1 0 R >>
```

Subsets are keyed by `fontCache[ v{0|1} | {fingerprint} | {baseName} | {runes} ]` so one deterministic subset per font per document is guaranteed, and hard caps bound untrusted input:

32 MiB reconstructed WOFF, 16 384 px image dimension, 16 Mpix pixel budget.

internal/pdf	Document model, operators, fonts, subsets, images, outlines (≈ 9.1k lines)
internal/pdf/assets	//go:embed Liberation + DejaVu faces; every accessor returns an isolated bytes.Clone
internal/convert	Caller: paints into Document, assembles TOC/outline/links/copies, then doc.Write
internal/imageout	Reuses fonts + shaping + glyph outlines; rasterizes instead of writing PDF bytes

- `Document` - object model; `AddPage`, `ReorderPages`, `DuplicatePage`, `SetOutline`, `Write`
- `Content.TextShow` - ASCII fast path, else shape + split mixed runs; records runes for the subsetter
- `ParseTTF` - sfnt validation; rejects CFF/OTTO and WOFF2 (Brotli not allowlisted)
- `ensureFont` - dispatch simple vs Type0; one subset per font per document (cache key above)
- `emitType0` - Identity-H CIDFontType2; CIDToGIDMap kept uncompressed for viewer safety
- `finalize` - ordered wiring: outlines before catalog, pages tree, then xref + trailer

Security: fonts are untrusted parse input — WOFF1 caps (tables ≤ 1024 , per-table ≤ 16 MiB, overlap rejection), bounds-checked tables, no remote https @font-face fetch, system fonts opt-in only. **Determinism:** byte-identical output for identical input; `CreationDate` is injectable via `Request.now()` and otherwise pins a fixed date.

display list → paint → content streams → xref (right edge, PDF mode) · deep-dive: [documentation/architecture/09-pdf-writer.md](#)

10

Image output & SVG raster

Image-mode tail: rasterize the shared display list to PNG/JPEG; bitmap fonts; SVG via `tdewolff/canvas`.

Image mode shares the whole *load* → *parse* → *style* → *layout* front half with PDF mode, then diverges at the paint/write tail: `internal/imageout` paints the layout **display list** into an in-memory `image.NRGBA` canvas and encodes a single PNG or JPEG. `internal/svg` rasterizes `` into PNG bytes through the allowlisted `tdewolff/canvas` path - the project's only third-party raster call.

display list → rasterize 2× NRGBA → box-filter → encode PNG/JPEG

Shared upstream, divergent sink - the Aug 2026 consolidation keeps text metrics, glyph forms and paint order on one table with PDF mode.

PDF mode

`internal/pdf` - Flate streams, font subsetting, xref

Image mode

`internal/imageout` - 2× supersample, no hinting, single page

`internal/imageout`

~3,000 lines: render options, supersampled rasterization, paint dispatch, decode/scale caches, PNG/JPEG encode

`internal/svg`

450 lines: `svg.Rasterize` via `tdewolff/canvas` only; `viewBox/width/height` sizing, panic containment

`internal/convert/render`

Shared `RenderObjects` → `Assemble` → `Finalize` lifecycle; `Assemble` is a no-op for image mode

`internal/pdf`

Consumed, not owned: `ShapeRun` text shaping, `GlyphContours`, `DefaultFont`, `ScanFontDirs`

- `RenderContext(ctx, root, opts)` — the core library render: layout result → rasterize → crop
- `RunRequest(ctx, req, log)` - engine seam shared by the `gowkhtmltoimage` CLI and `ImageConverter`
- `rasterizeContext` - paints `layout.Op` ops in `PaintOrder`, box-filters `rasterSS` supersample down
- `ttfDrawString` — TTF outline → bitmap glyphs (no FreeType, no hinting); fractional baseline, fake-bold pass
- `svg.Rasterize(data, maxSide)` — SVG → PNG bytes at intrinsic size; failure is a clean "no image"

Raster text uses **TTF outline AA**- AA via 2× supersampling with a box-filter downscale; no hinting. Fidelity claims are pinned in *10-imageout-svg.md* and *fidelity.md*; the shape was consolidated 2026-08-11.


Security & limits

- ACL merged (`Image.Load` $\oplus$ `Global.Load`); subresource fetches capped (64 fetches / 32 MiB)
- Decode hardening: ≤ 32 MiB encoded, dims $\leq 16,384$, ≤ 16 M pixels; canvas ≤ 64 M px / ≤ 256 MiB
- `--transparent` honored for PNG only; JPEG composites onto white with a warning
- Malformed SVG $\rightarrow$ clean error, never a crash (`defer recover()` in `rasterizeCanvas`)
- Single-page only: extra page objects and TOC are ignored with a warning
- Nearest-neighbour `<img>` scaling (Go 1.26 removed `image/draw` scalers)

Position: display list $\rightarrow$ raster $\rightarrow$ PNG/JPEG (right edge, image mode) · deep-dive: [documentation/architecture/10-imageout-svg.md](#)

Dependency DAG a strict downward graph toward the trust boundary

Nothing points back up. `internal/cli` is a leaf; the engine never parses argv. Library, image binary and PDF binary funnel into one neutral seam: `convert.Request`. `internal/load` sits at the bottom as the trust boundary - every subresource crosses the same two seams and inherits the same ACL, timeout and size caps.



`internal/load` - **trust boundary**: ACL, timeouts, redirects, size caps baked into control flow

- 1
- 2 `convert.Request` `render.Pipeline`
- 3 `cli.Command` `is`
- 4 `load.Loader.Load` `load.ResourceContext.Fetch`
- 5

PDF vs image mode same front half, different tail

<code>internal/pdf</code> - PDF 1.4	<code>internal/imageout</code> - PNG/JPEG
-------------------------------------	---

- content streams + zlib Flate, xref/trailer
- embedded Liberation Sans subsets, /Widths 1000-unit em
- Type0/CID for CJK, WOFF handling
- Catalog outlines, link annotations, metadata
- byte-stable output when `Now` is injected

- rasterize the shared display list - Assemble is a no-op
- bitmap font rasterizer (2× supersample, no hinting)
- single page; `--transparent` fill-alpha divergence
- SVG-as-image via `tdewo1ff/canvas` (only third-party raster call)
- reuses pdf shaping + outline data

Security architecture default-deny, enforced at the seams

ACL

- Local file access **denied by default**; opt-in via `--enable-local-file-access`
- `/ --allow` prefixes, symlinks resolved.

HTTP

- 30 s connect / 60 s response timeouts, max 10 redirects, 100 MiB body cap on both sides.

JS

- No JavaScript, ever - script content stripped at the layout stage.

BUDGETS

- Rule and resource budgets enforced at the convert layer, not scattered in callers.

HF

- Header/footer raw-markup rejected at the convert layer.

Full model: `documentation/THREAT-MODEL.md` · embedding guidance: `documentation/integration-security.md`.

Generated from `documentation/architecture/` - commit `ef526f9` (2026-08-11). Ten domain sections authored by ten parallel subagents; shell composed by the orchestrator. Pure HTML + CSS, no scripts.

MIT License · Copyright © 2026 Chinmay Sawant · [back to top ↑](#)